

SIMATS ENGINEERING

ASSESSMENT TOOL 1

COURSE NAME: DESIGN AND ANALYSIS OF ALGORITHMS

COURSE CODE: CSA06

NAME: A SHIVA SHANKAR REDDY

REG NO:192425174

COURSE OUTCOME COVERED

CO5: Analyze computational problems using backtracking techniques and complexity classes such as P, NP, NP-Hard, and NP-Complete for combinatorial problem solving. (BL5,BL6)

1. Problem Definition

A company has a set of jobs that must be assigned to a set of employees. Each job has a deadline by which it must be completed, and each employee can handle only one job at a time (no two jobs can be performed simultaneously by the same employee). The goal is to find a feasible assignment of jobs to employees and time slots such that every scheduled job is completed on or before its deadline, and no employee is double-booked at any point in time. Where it is not possible to schedule every job within the given resources, the goal becomes finding the maximum feasible subset of jobs that can be scheduled without violating any constraint.

This problem is a variant of the classical Job Sequencing with Deadlines problem, extended from a single machine/worker to multiple employees (parallel-machine scheduling). While the single-employee, single-slot-per-job version can be solved optimally using a greedy strategy, the multi-employee version with additional resource and requirement constraints is, in its general form, NP-Hard – equivalent to variants of parallel-machine scheduling and bin-packing problems for which no known polynomial-time algorithm exists. This makes Backtracking / Constraint Satisfaction (CSP) techniques a practical and correctness-guaranteeing choice.

1.1 Input

- A set of n jobs, each job J_i having a deadline d_i (the latest time slot by which it must be completed) and, optionally, a required skill/resource type and a processing duration.
- A set of m employees, each available to work one job per discrete time slot (1, 2, 3, ... up to the largest deadline among all jobs).

1.2 Expected Output

- A complete assignment mapping each job to an (employee, time-slot) pair such that all constraints are satisfied, if one exists.
- If a complete assignment is infeasible, the maximum feasible subset of jobs that can be scheduled without conflict, and a clear indication of which jobs could not be scheduled and why.

2. Identifying and Analyzing the Constraints

Constraint	Formal Description	Interaction / Effect on the Search
Time (Deadline) Constraint	Each job J_i must be assigned to a time slot t such that $1 \leq t \leq d_i$ (the job's deadline)	Restricts the domain of valid time slots for each job; jobs with smaller deadlines have fewer valid options and are more tightly constrained
Resource Availability Constraint	No employee may be assigned more than one job in the same time slot (all-different constraint on (employee, slot) pairs)	Creates direct competition between jobs whose deadlines force them toward the same early time slots
Job Requirement Constraint	A job may only be assigned to an employee who has the required skill/resource for that job (if specified)	Further restricts the domain of eligible employees for a given job, independent of time

These constraints interact multiplicatively rather than independently: a job's effective domain is the intersection of (a) all time slots up to its deadline, (b) all employees qualified for it, and (c)

all (employee, slot) pairs not already claimed by another job. Jobs with early deadlines and few qualified employees are the most tightly constrained and are therefore the most likely source of infeasibility – they should be prioritized first in the search, since if they cannot be placed, no ordering of the remaining jobs can rescue the schedule (this is the Most-Constrained-Variable / Minimum-Remaining-Values heuristic used in CSP search).

Constraint priority for the search is therefore: (1) Deadline constraint first – process jobs in increasing order of deadline, since their domains are smallest, (2) Resource availability constraint next – for a chosen job, try the earliest available (employee, slot) combination to leave the most flexibility for later jobs, (3) Job requirement (skill) constraint applied throughout as a domain filter, eliminating unqualified employee options before or during the availability check.

3. Formulating the Problem as a Constraint Satisfaction Problem (CSP)

CSP Element	Definition for This Problem
Variables	One variable per job: $J_1, J_2, \dots, J_n$
Domain of J_i	$\{ (e, t) : e \text{ is an employee qualified for } J_i, \text{ and } 1 \leq t \leq \text{deadline}(J_i) \}$
Constraints	All-Different constraint on assigned (e, t) pairs across all jobs (no employee double-booked in the same slot); each job's chosen slot must not exceed its deadline
Goal	Find a complete, consistent assignment of all variables (or the maximum consistent partial assignment if no complete one exists)

4. Backtracking / CSP Algorithm Design

The algorithm processes jobs one at a time in order of increasing deadline (Most-Constrained-Variable heuristic), and for each job tries every (employee, slot) pair in its domain, from the earliest available slot onward (Least-Constraining-Value heuristic, since using the earliest slot preserves the most flexibility for later jobs). If a chosen (employee, slot) pair is free, the job is assigned and the algorithm recurses to the next job. If the recursive call eventually fails to complete a valid full schedule, the current assignment is undone (backtracked) and the next candidate (employee, slot) pair is tried. If no candidate works for a job, the algorithm reports failure for that branch, causing the previous job in the ordering to backtrack and try an alternative assignment of its own.

5. Pseudocode

```

ALGORITHM ScheduleJobs(Jobs, Employees)
  INPUT : Jobs      = list of (jobID, deadline[, requiredSkill])
         Employees = list of employee IDs [with skill sets]
  OUTPUT: A feasible assignment {jobID -> (employee, slot)}, or the
         maximum feasible subset if a complete schedule is impossible

  SORT Jobs by deadline ASCENDING           // most-constrained-variable first
  used      <- empty set                     // set of (employee, slot) already
taken
  assignment <- empty map                    // jobID -> (employee, slot)
  unscheduled <- empty list

  CALL Backtrack(0)

```

```

RETURN assignment, unscheduled

FUNCTION Backtrack(index) RETURNS BOOLEAN
  IF index = LENGTH(Jobs) THEN
    RETURN TRUE // base case: every job placed
  END IF

  job <- Jobs[index]
  FOR t <- 1 TO job.deadline DO // earliest slot first (LCV
    heuristic)
    FOR EACH e IN Employees DO
      IF NOT Qualified(e, job) THEN // job-requirement constraint
        CONTINUE
      filter
    END IF
    IF (e, t) IN used THEN // resource-availability constraint
      CONTINUE
    END IF

    // ---- assign (tentative placement) ----
    used <- used UNION {(e, t)}
    assignment[job.id] <- (e, t)

    IF Backtrack(index + 1) THEN // success propagates back up
      RETURN TRUE
    END IF

    // ---- backtrack: undo the placement ----
    used <- used MINUS {(e, t)}
    REMOVE assignment[job.id]
  END FOR
END FOR

// No (employee, slot) worked for this job within its deadline
ADD job.id TO unscheduled
RETURN Backtrack(index + 1) // skip this job, continue with
rest
END FUNCTION

```

6. Pruning Strategies Applied

- Deadline-based domain restriction: a job's candidate slots are restricted to 1..deadline from the start, so the algorithm never even considers assigning a job to a slot that would violate its deadline – this is pruning by construction rather than by after-the-fact checking.
- Availability pruning: the 'used' set allows an $O(1)$ check to immediately skip any (employee, slot) combination already claimed, avoiding wasted recursive calls into branches that are guaranteed to fail the all-different constraint.
- Most-Constrained-Variable ordering: processing jobs with the smallest deadlines first means the tightest branches of the search tree are explored (and pruned, if infeasible) early, rather than discovering an unsolvable conflict only after many other jobs have already been placed.

- Least-Constraining-Value ordering: trying the earliest available slot first for each job leaves later slots open for subsequent, potentially more constrained jobs, reducing the likelihood that a full backtrack will be needed.
- Graceful degradation instead of total failure: if a job truly cannot be placed within its deadline given all other commitments, it is recorded as unscheduled and the search continues with the remaining jobs, rather than causing the entire algorithm to abandon an otherwise-good partial schedule.

7. Worked Example

7.1 Feasible Case – Full Schedule Found

Dataset: 5 jobs, 2 employees (E1, E2):

Job	Deadline
J1	2
J2	1
J3	2
J4	3
J5	1

Processing in order of increasing deadline (J2, J5, J1, J3, J4), the algorithm successfully places every job on its first attempt (no backtracking required in this instance), producing the complete, conflict-free schedule shown below:

```
J2 -> Employee E1, Slot 1
J5 -> Employee E2, Slot 1
J1 -> Employee E1, Slot 2
J3 -> Employee E2, Slot 2
J4 -> Employee E1, Slot 3
```

All 5 jobs are scheduled: every job's slot is $\leq$ its deadline, and no (employee, slot) pair is used more than once.

7.2 Infeasible Case – Backtracking in Action

Now add a sixth job, J6, with deadline 1, to the same dataset (2 employees). Three jobs (J2, J5, J6) all have deadline 1, but only 2 distinct (employee, slot) combinations exist for slot 1 (E1-slot1 and E2-slot1) – by the pigeonhole principle, all three cannot be scheduled simultaneously. Tracing the algorithm:

```
Assign J2 -> Employee E1, Slot 1
Assign J5 -> Employee E2, Slot 1
Backtrack: unassign J5 from Employee E2, Slot 1
Backtrack: unassign J2 from Employee E1, Slot 1
Assign J2 -> Employee E2, Slot 1
Assign J5 -> Employee E1, Slot 1
Backtrack: unassign J5 from Employee E1, Slot 1
Backtrack: unassign J2 from Employee E2, Slot 1
```

After exhausting every ordering of (employee, slot) combinations for the deadline-1 jobs, the algorithm correctly determines that J6 (the last deadline-1 job considered) cannot be placed, records it as unscheduled, and proceeds to schedule the remaining jobs (J1, J3, J4) normally. This demonstrates both the correctness of the backtracking mechanism (it exhaustively and

correctly proves infeasibility for the conflicting subset) and the value of graceful degradation (the remaining feasible jobs are still scheduled rather than the entire process failing).

8. Implementation (Python)

The following program implements the pseudocode of Section 5 and reproduces the worked examples of Section 7:

```
def schedule_jobs(jobs, employees):
    """jobs: list of (job_id, deadline). employees: list of employee ids."""
    jobs_sorted = sorted(jobs, key=lambda j: j[1])    # MCV: smallest deadline
    first
    n = len(jobs_sorted)
    used = set()                                     # (employee, slot) pairs already taken
    assignment = {}                                  # job_id -> (employee, slot)
    unscheduled = []

    def backtrack(i):
        if i == n:
            return True
        job_id, deadline = jobs_sorted[i]
        for t in range(1, deadline + 1):             # LCV: earliest slot first
            for e in employees:
                if (e, t) in used:                   # resource-availability
                    continue
                pruning
                used.add((e, t))
                assignment[job_id] = (e, t)
                if backtrack(i + 1):
                    return True
                used.remove((e, t))                  # backtrack: undo
                del assignment[job_id]
                unscheduled.append(job_id)            # graceful degradation
            return backtrack(i + 1)

    backtrack(0)
    return assignment, unscheduled

if __name__ == '__main__':
    employees = ['E1', 'E2']

    print('--- Feasible case (5 jobs) ---')
    jobs = [('J1', 2), ('J2', 1), ('J3', 2), ('J4', 3), ('J5', 1)]
    assignment, unscheduled = schedule_jobs(jobs, employees)
    for job_id, slot in sorted(assignment.items()):
        print(f'{job_id} -> Employee {slot[0]}, Slot {slot[1]}')
    print('Unscheduled:', unscheduled)

    print('\n--- Infeasible case (6 jobs, 3 share deadline 1) ---')
    jobs2 = jobs + [('J6', 1)]
    assignment2, unscheduled2 = schedule_jobs(jobs2, employees)
    for job_id, slot in sorted(assignment2.items()):
        print(f'{job_id} -> Employee {slot[0]}, Slot {slot[1]}')
    print('Unscheduled:', unscheduled2)
```


8.1 Program Output

```
--- Feasible case (5 jobs) ---
J1 -> Employee E1, Slot 2
J2 -> Employee E1, Slot 1
J3 -> Employee E2, Slot 2
J4 -> Employee E1, Slot 3
J5 -> Employee E2, Slot 1
Unscheduled: []
```

```
--- Infeasible case (6 jobs, 3 share deadline 1) ---
J1 -> Employee E1, Slot 2
J2 -> Employee E1, Slot 1
J3 -> Employee E2, Slot 2
J4 -> Employee E1, Slot 3
J5 -> Employee E2, Slot 1
Unscheduled: ['J6']
```

This output matches the manual trace in Sections 7.1 and 7.2, confirming the correctness of the algorithm.

OUTPUT:



The screenshot shows a code editor with a file named `main.py`. The code defines a `schedule_jobs` function that takes a list of jobs and a list of employees. It uses a backtracking algorithm to assign jobs to employees and slots. The function returns a tuple containing the assignment (a dictionary mapping job IDs to employee and slot) and a list of unscheduled jobs. The code also includes a `backtrack` function that recursively tries to assign jobs to employees and slots, pruning branches that are not feasible.

```
17 def schedule_jobs(jobs, employees):
18     """
19     jobs      : list of (job_id, deadline)
20     employees  : list of employee ids
21     returns   : (assignment: {job_id: (employee, slot)}, unscheduled:
22                  [job_id])
23     """
24     jobs_sorted = sorted(jobs, key=lambda j: j[1]) # MCV: smallest
25     # deadline first
26     n = len(jobs_sorted)
27     used = set() # (employee, slot) pairs already taken
28     assignment = {} # job_id -> (employee, slot)
29     unscheduled = []
30
31     def backtrack(i):
32         if i == n:
33             return True
34         job_id, deadline = jobs_sorted[i]
35         for t in range(1, deadline + 1): # LCV: earliest slot
36             first
37             for e in employees:
38                 if (e, t) in used:
39                     continue
40                 # resource-availability pruning
41                 used.add((e, t))
42                 assignment[job_id] = (e, t)
```

The output of the program is displayed on the right side of the editor. It shows the results for two test cases: a feasible case with 5 jobs and an infeasible case with 6 jobs. The output matches the manual trace in Sections 7.1 and 7.2.

```
--- Feasible case (5 jobs) ---
J1 -> Employee E1, Slot 2
J2 -> Employee E1, Slot 1
J3 -> Employee E2, Slot 2
J4 -> Employee E1, Slot 3
J5 -> Employee E2, Slot 1
Unscheduled: []

--- Infeasible case (6 jobs, 3 share deadline 1) ---
J1 -> Employee E1, Slot 2
J2 -> Employee E1, Slot 1
J3 -> Employee E2, Slot 2
J4 -> Employee E1, Slot 3
J5 -> Employee E2, Slot 1
Unscheduled: ['J6']

=== Code Execution Successful ===
```

9. Justification – How the Solution Satisfies All Constraints Efficiently

- Deadline constraint: satisfied by construction, since every candidate slot considered for a job is restricted to the range 1..deadline before any assignment is attempted – an out-of-range slot is never generated, let alone assigned.
- Resource availability constraint: satisfied by the 'used' set check, which guarantees in $O(1)$ time that no (employee, slot) pair is ever assigned to more than one job; this enforces the all-different constraint exactly and exhaustively.
- Job requirement constraint: satisfied by filtering the employee domain to only those qualified for a given job (`Qualified(e, job)` in the pseudocode) before checking availability, ensuring skill/resource compatibility is never violated.
- Efficiency: the Most-Constrained-Variable and Least-Constraining-Value heuristics reduce the size of the search tree actually explored, and infeasible branches are pruned immediately (rather than being fully expanded and checked only at the leaf), which is critical given the NP-Hard nature of the general problem – backtracking with these heuristics performs far better in practice than the theoretical worst case suggests, while still guaranteeing correctness (it will find a feasible complete schedule whenever one exists, and will correctly identify the maximum feasible subset otherwise).

10. Performance Analysis – Time and Space Complexity

Approach	Time Complexity	Space Complexity
Brute Force (try all assignments)	$O((m \times D)^n)$ where $D = \max$ deadline	$O(n)$
Backtracking with pruning + heuristics	Worst case exponential, e.g. $O((m \times D)^n)$ bound, but heavily reduced in practice by early pruning of infeasible (employee, slot) pairs	$O(n)$ for assignment/used-set + $O(n)$ recursion depth

Here n is the number of jobs, m is the number of employees, and D is the largest deadline among all jobs. Because the general multi-employee, deadline-constrained job scheduling problem is NP-Hard, no known algorithm can guarantee a polynomial-time solution for every instance; the exponential worst case of backtracking is therefore expected and unavoidable in the general case. However, the Most-Constrained-Variable ordering, Least-Constraining-Value slot selection, and immediate availability pruning used here mean that in most realistic instances – including both worked examples in Section 7 – the algorithm finds a solution (or correctly proves infeasibility for a subset) after examining only a small fraction of the theoretical search space, making it practical for typical company-scale job scheduling instances.

11. Suggested Tools and Platforms for Implementation

- Programming language: Python 3 (or Java/C++) to implement and test the backtracking/CSP algorithm.
- IDE / Editor: Visual Studio Code or PyCharm for development and debugging, especially useful for stepping through the recursive assignment/backtrack calls.
- Version control & submission: Git and GitHub, for hosting the Source_Code and Report as required by the submission instructions.
- CSP libraries (optional, for larger-scale problems): Google OR-Tools (CP-SAT solver) or python-constraint, which implement industrial-strength constraint propagation and search for large job-scheduling instances beyond what plain backtracking can handle efficiently.
- Testing: Manual test cases covering feasible schedules, infeasible/over-constrained schedules (as in Section 7.2), and edge cases (e.g., a job with a deadline smaller than any available slot, or an employee with no qualifying skills).

12. Conclusion

Job scheduling under deadline, resource-availability, and job-requirement constraints is naturally modeled as a Constraint Satisfaction Problem and solved using Backtracking, since the general multi-employee variant of the problem is NP-Hard and no efficient exact polynomial-time algorithm is known. By formulating jobs as variables, (employee, slot) pairs as domains, and the deadline/availability/requirement rules as constraints, and by applying Most-Constrained-Variable and Least-Constraining-Value heuristics together with immediate pruning of infeasible options, the algorithm designed and implemented above reliably finds a feasible complete schedule when one exists, and gracefully identifies the maximum feasible subset with a clear explanation of any unscheduled jobs when it does not. The Python implementation confirms the correctness of the pseudocode design through both a fully feasible worked example and a deliberately over-constrained example that exercises genuine backtracking behavior.